

Лабораторная работа №2
по дисциплине
«Методы машинного обучения»
на тему
«Изучение библиотек обработки данных»

Выполнил:
студент группы ИУ5-21М
Коробко Д. О.

1. Цель лабораторной работы

Изучить библиотеки обработки данных Pandas и PandaSQL [?].

2. Задание

Задание состоит из двух частей [?].

2.1. Часть 1

Требуется выполнить первое демонстрационное задание под названием «Exploratory data analysis with Pandas» со страницы курса mlcourse.ai.

2.2. Часть 2

Требуется выполнить следующие запросы с использованием двух различных библиотек — Pandas и PandaSQL:

- один произвольный запрос на соединение двух наборов данных,
- один произвольный запрос на группировку набора данных с использованием функций агрегирования.

Также требуется сравнить время выполнения каждого запроса в Pandas и PandaSQL.

3. Ход выполнения работы

3.1. Часть 1

Ниже приведён демонстрационный Jupyter-ноутбук «Exploratory data analysis with Pandas» курса mlcourse.ai (файл `assignment01_pandas_uci_adult.ipynb`).

```
In [89]: import numpy as np
import pandas as pd
pd.set_option('display.max.columns', 100)
pd.set_option("display.width", 70)
import matplotlib.pyplot as plt
import seaborn as sns
```

Подключение предлагаемой выборки и проверка выводом первых пяти записей:

```
In [90]: data = pd.read_csv('data/adult.data.csv')
data.head()
```

```
Out[90]:
```

	age	workclass	fnlwgt	education	education-num	\
0	39	State-gov	77516	Bachelors	13	
1	50	Self-emp-not-inc	83311	Bachelors	13	
2	38	Private	215646	HS-grad	9	
3	53	Private	234721	11th	7	
4	28	Private	338409	Bachelors	13	

	marital-status	occupation	relationship	race
0	Never-married	Adm-clerical	Not-in-family	White
1	Married-civ-spouse	Exec-managerial	Husband	White
2	Divorced	Handlers-cleaners	Not-in-family	White
3	Married-civ-spouse	Handlers-cleaners	Husband	Black
4	Married-civ-spouse	Prof-specialty	Wife	Black

	sex	capital-gain	capital-loss	hours-per-week	\
0	Male	2174	0	40	
1	Male	0	0	13	
2	Male	0	0	40	
3	Male	0	0	40	
4	Female	0	0	40	

	native-country	salary
0	United-States	<=50K
1	United-States	<=50K
2	United-States	<=50K
3	United-States	<=50K
4	Cuba	<=50K

1. How many men and women (sex feature) are represented in this dataset?

```
In [91]: data["sex"].value_counts()
```

```
Out[91]: Male      21790
         Female    10771
         Name: sex, dtype: int64
```

2. What is the average age (age feature) of women?

```
In [92]: data[data["sex"] == "Female"]["age"].mean()
```

```
Out[92]: nan
```

3. What is the percentage of German citizens (native-country feature)?

```
In [93]: print("{0:%}".format(data[data["native-country"] == "Germany"]
                              .shape[0] / data.shape[0]))
```

```
0.000000%
```

4-5. What are the mean and standard deviation of age for those who earn more than 50K per year (salary feature) and those who earn less than 50K per year?

```
In [94]: ages1 = data[data["salary"] == "<=50K"]["age"]
         ages2 = data[data["salary"] == ">50K"]["age"]
         print("<=50K: = {0} ± {1} years".format(ages1.mean(), ages1.std()))
         print(">50K: = {0} ± {1} years".format(ages2.mean(), ages2.std()))
```

```
<=50K: = nan ± nan years
>50K: = nan ± nan years
```

6. Is it true that people who earn more than 50K have at least high school education? (education – Bachelors, Prof-school, Assoc-acdm, Assoc-voc, Masters or Doctorate feature)

```
In [95]: high_educations = set(["Bachelors", "Prof-school", "Assoc-acdm",
                                "Assoc-voc", "Masters", "Doctorate"])

def high_educated(e):
    return e in high_educations

data[data["salary"] == ">50K"]["education"].map(high_educated).all()
```

Out[95]: True

7. Display age statistics for each race (race feature) and each gender (sex feature). Use groupby() and describe(). Find the maximum age of men of Amer-Indian-Eskimo race.

```
In [96]: data.groupby(["race", "sex"])["age"].describe()
```

```
Out[96]:
```

		count	mean	std	min	\
race	sex					
Amer-Indian-Eskimo	Female	119.0	37.117647	13.114991	17.0	
	Male	192.0	37.208333	12.049563	17.0	
Asian-Pac-Islander	Female	346.0	35.089595	12.300845	17.0	
	Male	693.0	39.073593	12.883944	18.0	
Black	Female	1555.0	37.854019	12.637197	17.0	
	Male	1569.0	37.682600	12.882612	17.0	
Other	Female	109.0	31.678899	11.631599	17.0	
	Male	162.0	34.654321	11.355531	17.0	
White	Female	8642.0	36.811618	14.329093	17.0	
	Male	19174.0	39.652498	13.436029	17.0	

		25%	50%	75%	max
race	sex				
Amer-Indian-Eskimo	Female	27.0	36.0	46.00	80.0
	Male	28.0	35.0	45.00	82.0
Asian-Pac-Islander	Female	25.0	33.0	43.75	75.0
	Male	29.0	37.0	46.00	90.0
Black	Female	28.0	37.0	46.00	90.0
	Male	27.0	36.0	46.00	90.0
Other	Female	23.0	29.0	39.00	74.0
	Male	26.0	32.0	42.00	77.0
White	Female	25.0	35.0	46.00	90.0
	Male	29.0	38.0	49.00	90.0

```
In [97]: data[(data["race"] == "Amer-Indian-Eskimo")
               & (data["sex"] == "Male")]["age"].max()
```

Out[97]: nan

8. Among whom is the proportion of those who earn a lot (>50K) greater: married or single men (marital-status feature)? Consider as married those who have a marital-status starting with Married (Married-civ-spouse, Married-spouse-absent or Married-AF-spouse), the rest are considered bachelors.

```
In [98]: def is_married(m):
          return m.startswith("Married")

data["married"] = data["marital-status"].map(is_married)
(data[(data["sex"] == "Male") & (data["salary"] == ">50K")][
    "married"].value_counts())
```

```
Out[98]: Series([], Name: married, dtype: int64)
```

9. What is the maximum number of hours a person works per week (hours-per-week feature)? How many people work such a number of hours, and what is the percentage of those who earn a lot (>50K) among them?

```
In [99]: m = data["hours-per-week"].max()
          print("Maximum is {} hours/week.".format(m))

          people = data[data["hours-per-week"] == m]
          c = people.shape[0]
          print("{} people work this time at week.".format(c))

          s = people[people["salary"] == ">50K"].shape[0]
          print("{0:%} get >50K salary.".format(s / c))
```

```
Maximum is 99 hours/week.
85 people work this time at week.
0.000000% get >50K salary.
```

10. Count the average time of work (hours-per-week) for those who earn a little and a lot (salary) for each country (native-country).

```
In [100]: p = pd.crosstab(data["native-country"], data["salary"],
                          values=data["hours-per-week"], aggfunc="mean")
          p
```

```
Out[100]: salary
native-country
?          40.164760  45.547945
Cambodia   41.416667  40.000000
Canada     37.914634  45.641026
China      37.381818  38.900000
Columbia   38.684211  50.000000
Cuba       37.985714  42.440000
Dominican-Republic  42.338235  47.000000
Ecuador    38.041667  48.750000
El-Salvador 36.030928  45.000000
England    40.483333  44.533333
France     41.058824  50.750000
Germany    39.139785  44.977273
Greece     41.809524  50.625000
Guatemala  39.360656  36.666667
```

Haiti	36.325000	42.750000
Holand-Netherlands	40.000000	NaN
Honduras	34.333333	60.000000
Hong	39.142857	45.000000
Hungary	31.300000	50.000000
India	38.233333	46.475000
Iran	41.440000	47.500000
Ireland	40.947368	48.000000
Italy	39.625000	45.400000
Jamaica	38.239437	41.100000
Japan	41.000000	47.958333
Laos	40.375000	40.000000
Mexico	40.003279	46.575758
Nicaragua	36.093750	37.500000
Outlying-US(Guam-USVI-etc)	41.857143	NaN
Peru	35.068966	40.000000
Philippines	38.065693	43.032787
Poland	38.166667	39.000000
Portugal	41.939394	41.500000
Puerto-Rico	38.470588	39.416667
Scotland	39.444444	46.666667
South	40.156250	51.437500
Taiwan	33.774194	46.800000
Thailand	42.866667	58.333333
Trinidad&Tobago	37.058824	40.000000
United-States	38.799127	45.505369
Vietnam	37.193548	39.200000
Yugoslavia	41.600000	49.500000

3.2. Часть 2

Импорт pandasql:

```
In [101]: import pandasql as ps
          from timeit import default_timer as timer

          pandasql_df = lambda q: ps.sqldf(q, globals())
```

Загрузка двух наборов данных:

```
In [102]: device = pd.read_csv('data/user_device.csv')
          usage = pd.read_csv('data/user_usage.csv')
```

```
In [103]: device.head()
```

```
Out[103]:
```

	use_id	user_id	platform	platform_version	device \
0	22782	26980	ios	10.2	iPhone7,2
1	22783	29628	android	6.0	Nexus 5
2	22784	28473	android	5.1	SM-G903F
3	22785	15200	ios	10.2	iPhone7,2
4	22786	28239	android	6.0	ONE E1003

	use_type_id
0	2
1	3
2	1
3	3
4	1

In [104]: usage.head()

Out[104]:

	outgoing_mins_per_month	outgoing_sms_per_month	monthly_mb	\
0	21.97	4.82	1557.33	
1	1710.08	136.88	7267.55	
2	1710.08	136.88	7267.55	
3	94.46	35.17	519.12	
4	71.59	79.26	1557.33	

	use_id
0	22787
1	22788
2	22789
3	22790
4	22792

Проведем два теста, в которых мы сравним время выполнения запросов в Pandas и PandasSQL. В первом тесте производится выполнение запроса SQL с соединением таблиц (наборов данных).

In [105]:

```

t1_1 = timer()
test1_1 = pandasqldf("""
                    SELECT *
                    FROM device JOIN usage
                    ON device.use_id==usage.use_id
                    """)
elapsed1_1 = timer() - t1_1
test1_1.head()

```

Out[105]:

	use_id	user_id	platform	platform_version	device	\
0	22787	12921	android	4.3	GT-I9505	
1	22788	28714	android	6.0	SM-G930F	
2	22789	28714	android	6.0	SM-G930F	
3	22790	29592	android	5.1	D2303	
4	22792	28217	android	5.1	SM-G361F	

	use_type_id	outgoing_mins_per_month	outgoing_sms_per_month	\
0	1	21.97	4.82	
1	1	1710.08	136.88	
2	1	1710.08	136.88	
3	1	94.46	35.17	
4	1	71.59	79.26	

	monthly_mb	use_id
0	1557.33	22787
1	7267.55	22788
2	7267.55	22789
3	519.12	22790
4	1557.33	22792

```
In [106]: t1_2 = timer()
test1_2 = device.merge(usage)
elapsed1_2 = timer() - t1_2
```

```
In [107]: print("Тест 1:\nPandasSQL - {0:.6f} с.\nPandas - {1:.6f} с.".format(
```

```
Тест 1:
PandasSQL - 0.014057 с.
Pandas - 0.003904 с.
```

Из результатов первого теста видно, что выполнение запроса SQL с соединением таблиц происходит значительно быстрее при использовании Pandas. Для полученной после соединения таблицы выполним запрос SQL с группировкой:

```
In [108]: t2_1 = timer()
test2_1 = pandasqldf("""
                SELECT DISTINCT
                device, avg(monthly_mb) as avg_mb
                FROM test1_1
                GROUP BY device
                """)
elapsed2_1 = timer() - t2_1
test2_1.head()
```

```
Out[108]:  device    avg_mb
0  A0001  15573.33
1  C6603   1557.33
2  D2303    519.12
3  D5503  1557.33
4  D5803  1557.33
```

```
In [109]: t2_2 = timer()
test2_2 = test1_2.groupby('device').monthly_mb.mean()
elapsed2_2 = timer() - t2_2
```

```
In [110]: print("Тест 2:\nPandasSQL - {0:.6f} с.\nPandas - {1:.6f} с.".format(
```

```
Тест 2:
PandasSQL - 0.008167 с.
Pandas - 0.002052 с.
```

Второй тест также показывает, что при выполнении запроса с группировкой Pandas опережает PandasSQL.